

5

Sorting pairwise sums

Introduction

Let A be some linearly ordered set and $(\oplus) :: A \rightarrow A \rightarrow A$ some monotonic binary operation on A , so $x \leq x' \wedge y \leq y' \Rightarrow x \oplus y \leq x' \oplus y'$. Consider the problem of computing

$$\begin{aligned} \text{sortsums} &:: [A] \rightarrow [A] \rightarrow [A] \\ \text{sortsums } xs \ ys &= \text{sort } [x \oplus y \mid x \leftarrow xs, y \leftarrow ys] \end{aligned}$$

Counting just comparisons, and supposing xs and ys have the same length n , how long does $\text{sortsums } xs \ ys$ take?

Certainly $O(n^2 \log n)$ comparisons are sufficient. There are n^2 sums and sorting a list of length n^2 can be done with $O(n^2 \log n)$ comparisons. This upper bound does not depend on \oplus being monotonic. In fact, without further information about \oplus and A this bound is also a lower bound. The assumption that \oplus is monotonic does not reduce the asymptotic complexity, only the constant factor.

But now suppose we know more about \oplus and A : specifically that (\oplus, A) is an *Abelian group*. Thus, \oplus is associative and commutative, with identity element e and an operation $\text{negate} :: A \rightarrow A$ such that $x \oplus \text{negate } x = e$. Given this extra information, [Jean-Luc Lambert \(1992\)](#) proved that sortsums can be computed with $O(n^2)$ comparisons. However, his algorithm also requires $Cn^2 \log n$ additional operations, where C is quite large. It remains an open problem, some 35 years after it was first posed by [Harper et al. \(1975\)](#), as to whether the total cost of computing sortsums can be reduced to $O(n^2)$ comparisons and $O(n^2)$ other steps.

Lambert's algorithm is another nifty example of divide and conquer. Our aim in this pearl is just to present the essential ideas and give an implementation in Haskell.

Lambert's algorithm

Let's first prove the $\Omega(n^2 \log n)$ lower bound on *sortsums* when the only assumption is that (\oplus) is monotonic. Suppose *xs* and *ys* are both sorted into increasing order and consider the $n \times n$ matrix

$$[[x \oplus y \mid y \leftarrow ys] \mid x \leftarrow xs]$$

Each row and column of the matrix is therefore in increasing order. The matrix is an example of a standard Young tableau, and it follows from Theorem H of Section 5.1.4 of [Knuth \(1998\)](#) that there are precisely

$$E(n) = (n^2)! \left/ \left(\frac{(2n-1)!}{(n-1)!} \frac{(2n-2)!}{(n-2)!} \cdots \frac{n!}{0!} \right) \right.$$

ways of assigning the values 1 to n^2 to the elements of the matrix, and so exactly $E(n)$ potential permutations that sort the input. Using the fact that $\log E(n) = \Omega(n^2 \log n)$, we conclude that at least this number of comparisons is required.

Now for the meat of the exercise. Lambert's algorithm depends on two simple facts. Define the subtraction operation $(\ominus) :: A \rightarrow A \rightarrow A$ by $x \ominus y = x \oplus \text{negate } y$. Then:

$$x \oplus y = x \ominus \text{negate } y \tag{5.1}$$

$$x \ominus y \leq x' \ominus y' \equiv x \ominus x' \leq y \ominus y' \tag{5.2}$$

Verification of (5.1) is easy, but (5.2), which we leave as an exercise, requires all the properties of an Abelian group. In effect, (5.1) says that the problem of sorting sums can be reduced to the problem of sorting subtractions and (5.2) says that the latter problem is, in turn, reducible to the problem of sorting subtractions over a single list.

Here is how (5.1) and (5.2) are used. Consider the list *subs xs ys* of labelled subtractions defined by

$$\begin{aligned} \text{subs} &:: [A] \rightarrow [A] \rightarrow [\text{Label } A] \\ \text{subs } xs \text{ } ys &= [(x \ominus y, (i, j)) \mid (x, i) \leftarrow \text{zip } xs [1..], (y, j) \leftarrow \text{zip } ys [1..]] \end{aligned}$$

where *Label a* is a synonym for $(a, (Int, Int))$. Thus, each term $x \ominus y$ is labelled with the position of *x* in *xs* and *y* in *ys*. Labelling information will be needed later on. The first fact (5.1) gives

$$\begin{aligned} \text{sortsums } xs \text{ } ys &= \text{map fst } (\text{sortsubs } xs \text{ } (\text{map negate } ys)) \\ \text{sortsubs } xs \text{ } ys &= \text{sort } (\text{subs } xs \text{ } ys) \end{aligned}$$

The sums are sorted by sorting the associated labelled subtractions and throwing away the labels.


```

sortsums xs ys = map fst (sortsubs xs (map negate ys))
sortsubs xs ys = sortBy (cmp (mkArray xs ys)) (subs xs ys)

subs xs ys = [(x ⊖ y, (i, j)) | (x, i) ← zip xs [1..], (y, j) ← zip ys [1..]]

cmp a (x, (i, j)) (y, (k, ℓ)) = compare (a ! (1, i, k)) (a ! (2, j, ℓ))

mkArray xs ys = array b (zip (table xs ys) [1..])
                where b = ((1, 1, 1), (2, p, p))
                      p = max (length xs) (length ys)

table xs ys = map snd (map (tag 1) xxs ∧ map (tag 2) yys)
                where xxs = sortsubs' xs
                      yys = sortsubs' ys

tag i (x, (j, k)) = (x, (i, j, k))

```

Fig. 5.1 The complete code for *sortsums*, except for *sortsubs'*

well founded. Although computing *sortsubs xs ys* takes $O(mn \log mn)$ steps, it uses no comparisons on *A* beyond those needed to construct *table*. And *table* needs only $O(m^2 + n^2)$ comparisons plus those comparisons needed to construct *sortsubs' xs* and *sortsubs' ys*. What remains is to show how to compute *sortsubs'* with a quadratic number of comparisons.

Divide and conquer

Ignoring labels for the moment and writing $xs \ominus ys$ for $[x \ominus y \mid x \leftarrow xs, y \leftarrow ys]$, the key to a divide and conquer algorithm is the identity

$$\begin{aligned}
 (xs \uplus ys) \ominus (xs \uplus ys) \\
 = (xs \ominus xs) \uplus (xs \ominus ys) \uplus (ys \ominus xs) \uplus (ys \ominus ys)
 \end{aligned}$$

Hence, to sort the list on the left, we can sort the four lists on the right and merge them together. The presence of labels complicates the divide and conquer algorithm slightly because the labels have to be adjusted correctly. The labelled version reads

$$\begin{aligned}
 \text{subs } (xs \uplus ys) (xs \uplus ys) \\
 = \text{subs } xs \, xs \uplus \text{map } (\text{incr } m) (\text{subs } xs \, ys) \uplus \\
 \text{map } (\text{incl } m) (\text{subs } ys \, xs) \uplus \text{map } (\text{incb } m) (\text{subs } ys \, ys)
 \end{aligned}$$

where $m = \text{length } xs$ and

$$\begin{aligned}
 \text{incl } m (x, (i, j)) &= (x, (m+i, j)) \\
 \text{incr } m (x, (i, j)) &= (x, (i, m+j)) \\
 \text{incb } m (x, (i, j)) &= (x, (m+i, m+j))
 \end{aligned}$$

```

sortsubs' [] = []
sortsubs' [w] = [(w ⊖ w, (1, 1))]
sortsubs' ws = foldr1 (⋈) [xs, map (incr m) xys,
                          map (incl m) yxs, map (incb m) yys]
  where xs = sortsubs' xs
        xys = sortBy (cmp (mkArray xs ys)) (subs xs ys)
        yxs = map switch (reverse xys)
        yys = sortsubs' ys
        (xs, ys) = splitAt m ws
        m = length ws div 2

incl m (x, (i, j)) = (x, (m + i, j))
incr m (x, (i, j)) = (x, (i, m + j))
incb m (x, (i, j)) = (x, (m + i, m + j))

switch (x, (i, j)) = (negate x, (j, i))

```

Fig. 5.2 The code for *sortsubs'*

To compute *sortsubs' ws* we split *ws* into two equal halves *xs* and *ys*. The lists *sortsubs' xs* and *sortsubs' ys* are computed recursively. The list *sortsubs xs ys* is computed by applying the algorithm of the previous section. We can also compute *sortsubs ys xs* in the same way, but an alternative is simply to reverse *sortsubs xs ys* and negate its elements:

```

sortsubs ys xs = map switch (reverse (sortsubs xs ys))
switch (x, (i, j)) = (negate x, (j, i))

```

The program for *sortsubs'* is given in [Figure 5.2](#). The number $C(n)$ of comparisons required to compute *sortsubs'* on a list of length n satisfies the recurrence $C(n) = 2C(n/2) + O(n^2)$ with solution $C(n) = O(n^2)$. That means *sortsums* can be computed with $O(n^2)$ comparisons. However, the total time $T(n)$ satisfies $T(n) = 2T(n/2) + O(n^2 \log n)$ with solution $T(n) = O(n^2 \log n)$. The logarithmic factor can be removed from $T(n)$ if *sortBy cmp* can be computed in quadratic time, but this result remains elusive. In any case, the additional complexity arising from replacing comparisons by other operations makes the algorithm very inefficient in practice.

Final remarks

The problem of sorting pairwise sums is given as Problem 41 in the Open Problems Project ([Demaine et al., 2009](#)), a web resource devoted to recording open problems of interest to researchers in computational geometry and related fields. The earliest known reference to the problem is [Fedman \(1976\)](#),

who attributes the problem to Elwyn Berlekamp. All these references consider the problem in terms of numbers rather than Abelian groups, but the idea is the same.

References

- Demaine, E. D., Mitchell, J. S. B. and O'Rourke, J. (2009). The Open Problems Project. <http://mave.smith.edu/~orourke/TOPP/>.
- Fedman, M. L. (1976). How good is the information theory lower bound in sorting? *Theoretical Computer Science* **1**, 355–61.
- Harper, L. H., Payne, T. H., Savage, J. E. and Straus, E. (1975). Sorting $X + Y$. *Communications of the ACM* **18** (6), 347–9.
- Knuth, D. E. (1998). *The Art of Computer Programming: Volume 3, Sorting and Searching*, second edition. Reading, MA: Addison-Wesley.
- Lambert, J.-L. (1992). Sorting the sums $(x_i + y_j)$ in $O(n^2)$ comparisons. *Theoretical Computer Science* **103**, 137–41.